

```
In [1]: from sklearn.datasets import load_breast_cancer
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = load_breast_cancer(as_frame=True)
```

```
In [3]: X = df.data
y = df.target
```

```
In [4]: X
```

```
Out[4]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14710
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.08690	0.07060
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.19740	0.12730
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.24140	0.10520
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.19800	0.10430
...
564	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13860
565	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09710
566	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	0.05360
567	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.35140	0.15260
568	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00000	0.00000

569 rows x 30 columns

```
In [5]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
import numpy as np
from sklearn.model_selection import GridSearchCV
```

```
In [6]: knn_pipe = Pipeline([
    ("scaler", StandardScaler()),
    ("knn", KNeighborsClassifier())
])
```

```
In [7]: knn_grid = {  
        "knn__n_neighbors": np.arange(1,31,2)  
    }
```

```
In [8]: knn = GridSearchCV(  
        knn_pipe,  
        knn_grid,  
        cv = 5,  
        n_jobs = -1  
    )
```

```
In [9]: lr_pipe = Pipeline([  
        ("scaler", StandardScaler()),  
        ("lr", LogisticRegression(max_iter=200_000))  
    ])
```

```
In [10]: lr_grid = {  
        "lr__C": np.logspace(-3, 3, 50)  
    }
```

```
In [11]: lr = GridSearchCV(  
        lr_pipe,  
        lr_grid,  
        cv = 5,  
        n_jobs = -1  
    )
```

```
In [12]: dt_pipe = Pipeline([  
        ("scaler", StandardScaler()),  
        ("dt", DecisionTreeClassifier())  
    ])
```

```
In [13]: DecisionTreeClassifier()
```

```
Out[13]: ▼ DecisionTreeClassifier ⓘ ?
```

► Parameters

```
In [14]: dt_grid = {  
        "dt__criterion": ["gini", "entropy"],  
        "dt__max_depth": [None, 3, 5, 8, 12]  
    }
```

```
In [15]: dt = GridSearchCV(  
        dt_pipe,  
        dt_grid,  
        cv = 5,  
        n_jobs = -1  
    )
```

```
In [16]: nb_pipe = Pipeline([
    ("scaler", StandardScaler()),
    ("nb", GaussianNB())
])
```

```
In [17]: nb_grid = {
    "nb__var_smoothing": np.logspace(-10, -8, 20)
}
```

```
In [18]: nb = GridSearchCV(
    nb_pipe,
    nb_grid,
    cv = 5,
    n_jobs = -1
)
```

```
In [19]: rf_pipe = Pipeline([
    ("scaler", StandardScaler()),
    ("rf", RandomForestClassifier())
])
```

```
In [20]: RandomForestClassifier()
```

```
Out[20]: ▼ RandomForestClassifier ⓘ ?
```

```
► Parameters
```

```
In [21]: rf_grid = {
    "rf__n_estimators": [50, 100, 200],
    "rf__max_depth": [5, 10, 15, None],
    "rf__min_samples_split": [2, 5, 10]
}
```

```
In [22]: rf = GridSearchCV(
    rf_pipe,
    rf_grid,
    cv = 5,
    n_jobs = -1
)
```

```
In [23]: from sklearn.model_selection import StratifiedKFold
import pandas as pd
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
    f1_score, confusion_matrix, roc_curve, auc
)
import matplotlib.pyplot as plt
```

```
In [24]: def cv5_report(name, model, X, y, outer_splits=5):

    skf = StratifiedKFold(n_splits = outer_splits, random_state = 42, shuffle

    folds = []
```

```

aucs = []

plt.figure(figsize=(6,5))
for fold, (train_index, test_index) in enumerate(skf.split(X,y),1):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    model.fit(X_train, y_train)

    if hasattr(model, "best_estimator_"):
        best_model = model.best_estimator_
    else:
        best_model = model

    y_test_pred = best_model.predict(X_test)
    y_test_prob = best_model.predict_proba(X_test)[:,1]

    #Test Metrics
    acc = accuracy_score(y_test, y_test_pred)
    pre = precision_score(y_test, y_test_pred)
    rec = recall_score(y_test, y_test_pred)
    f1 = f1_score(y_test, y_test_pred)

    #ROC_curve
    fpr, tpr, _ = roc_curve(y_test, y_test_prob)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f"AUC = {roc_auc:.2f}")
    plt.xlabel("False Positive Rate (FPR)")
    plt.ylabel("True Positive Rate (TPR)")
    plt.title(f"{name} ROC Curve")
    plt.grid()
    plt.legend()

    aucs.append(roc_auc)

    folds.append(
        {"Model":name, "Fold":fold, "Accuracy":acc, "Precision":pre,
        "Recall":rec, "F1":f1, "AUC": roc_auc}
    )

plt.show()

folds_df = pd.DataFrame(folds)

row_mean = {
    "Model":name,
    "Fold":"Mean",
    "Accuracy":folds_df["Accuracy"].mean(),
    "Precision":folds_df["Precision"].mean(),
    "Recall":folds_df["Recall"].mean(),
    "F1":folds_df["F1"].mean(),
    "AUC":folds_df["AUC"].mean(),
}

results = pd.concat([folds_df, pd.DataFrame([row_mean])])

```

```
return results
```



```
In [25]: from sklearn.svm import SVC
```



```
In [26]: svm_pipe = Pipeline([
    ("scaler", StandardScaler()),
    ("svm", SVC(probability = True))
])
```


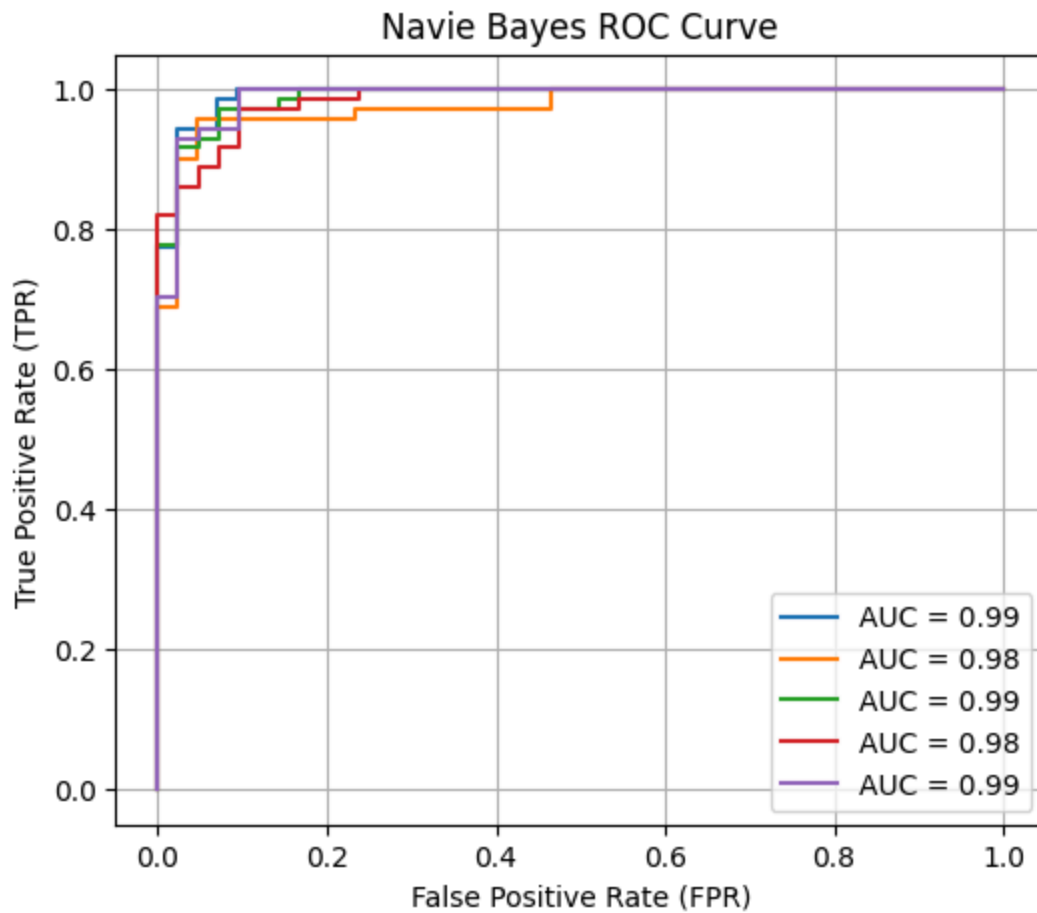
```
In [27]: svm_grid = {
    "svm__C": np.logspace(-2,2,50),
    "svm__kernel" : ["linear"]
}
```

```
In [28]: svm = GridSearchCV(
    svm_pipe,
    svm_grid,
    cv = 5,
    n_jobs = -1
)
```

```
In [29]: pd.concat([
    cv5_report("LR", lr, X, y, 5),
    cv5_report("KNN", knn, X, y, 5),
    cv5_report("SVM", svm, X, y, 5),
    cv5_report("DecisionTree", dt, X, y, 5),
    cv5_report("Navie Bayes", nb, X, y, 5),
    cv5_report("Random Forest", rf, X, y, 5)
]).set_index(["Model"])
```







Out [29]:

	Fold	Accuracy	Precision	Recall	F1	AUC
Model						
LR	1	0.973684	0.972222	0.985915	0.979021	0.998362
LR	2	0.938596	0.910256	1.000000	0.953020	0.997707
LR	3	0.964912	0.947368	1.000000	0.972973	0.985780
LR	4	0.991228	1.000000	0.986111	0.993007	1.000000
LR	5	0.991150	0.986111	1.000000	0.993007	0.998659
LR	Mean	0.971914	0.963192	0.994405	0.978206	0.996102
KNN	1	0.982456	0.972603	1.000000	0.986111	0.997380
KNN	2	0.947368	0.933333	0.985915	0.958904	0.973305
KNN	3	0.947368	0.934211	0.986111	0.959459	0.975694
KNN	4	0.982456	1.000000	0.972222	0.985915	0.999339
KNN	5	0.973451	0.959459	1.000000	0.979310	0.984574
KNN	Mean	0.966620	0.959921	0.988850	0.973940	0.986058
SVM	1	0.982456	0.985915	0.985915	0.985915	0.998035
SVM	2	0.956140	0.945946	0.985915	0.965517	0.996397
SVM	3	0.964912	0.947368	1.000000	0.972973	0.985119
SVM	4	0.991228	1.000000	0.986111	0.993007	0.999339
SVM	5	0.991150	0.986111	1.000000	0.993007	0.998659
SVM	Mean	0.977177	0.973068	0.991588	0.982084	0.995510
DecisionTree	1	0.938596	0.944444	0.957746	0.951049	0.931707
DecisionTree	2	0.921053	0.918919	0.957746	0.937931	0.909106
DecisionTree	3	0.912281	0.887500	0.986111	0.934211	0.911541
DecisionTree	4	0.947368	0.985294	0.930556	0.957143	0.953373
DecisionTree	5	0.920354	0.897436	0.985915	0.939597	0.949866
DecisionTree	Mean	0.927930	0.926719	0.963615	0.943986	0.931118
Navie Bayes	1	0.956140	0.958333	0.971831	0.965035	0.991811
Navie Bayes	2	0.912281	0.906667	0.957746	0.931507	0.976089
Navie Bayes	3	0.929825	0.910256	0.986111	0.946667	0.988757
Navie Bayes	4	0.903509	0.969231	0.875000	0.919708	0.984788
Navie Bayes	5	0.946903	0.945205	0.971831	0.958333	0.988598
Navie Bayes	Mean	0.929731	0.937939	0.952504	0.944250	0.986009
Random Forest	1	0.964912	0.985507	0.957746	0.971429	0.998035

	Fold	Accuracy	Precision	Recall	F1	AUC
Model						
Random Forest	2	0.921053	0.897436	0.985915	0.939597	0.978873
Random Forest	3	0.973684	0.960000	1.000000	0.979592	0.981647
Random Forest	4	0.947368	0.985294	0.930556	0.957143	0.996032
Random Forest	5	0.964602	0.946667	1.000000	0.972603	0.991449
Random Forest	Mean	0.954324	0.954981	0.974844	0.964073	0.989207

In [30]: `knn.best_params_`

Out[30]: `{'knn__n_neighbors': np.int64(7)}`

In [31]: `lr.best_params_`

Out[31]: `{'lr__C': np.float64(0.655128556859551)}`

In [32]: `svm.best_params_`

Out[32]: `{'svm__C': np.float64(1.325711365590108), 'svm__kernel': 'linear'}`

In [33]: `nb.best_params_`

Out[33]: `{'nb__var_smoothing': np.float64(1e-10)}`

In [34]: `rf.best_params_`

Out[34]: `{'rf__max_depth': 15, 'rf__min_samples_split': 2, 'rf__n_estimators': 50}`

```
In [35]: # Get feature importances from the best model

best_rf = rf.best_estimator_
feature_importance_rf = best_rf.named_steps["rf"].feature_importances_

# Create a DataFrame to visualize feature importance

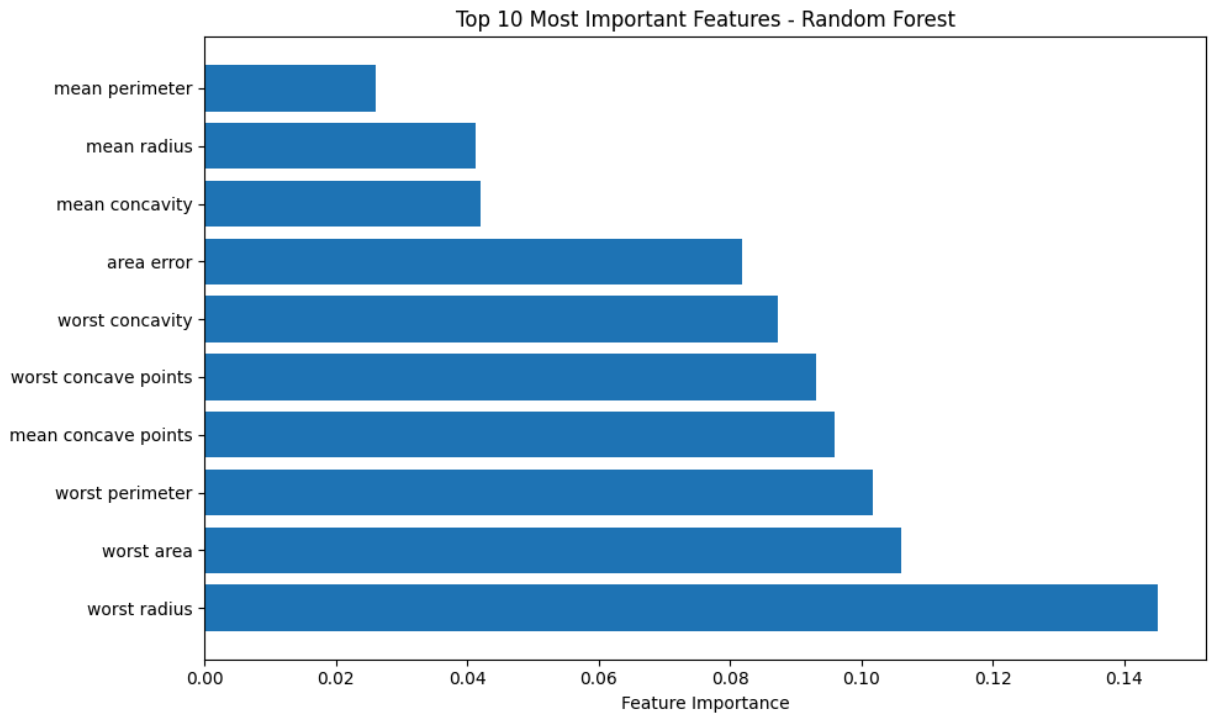
importance_df_rf = pd.DataFrame({
    'Feature': X.columns,
    'Importance': feature_importance_rf
}).sort_values('Importance', ascending=False)
print("\nTop 10 Most Important Features - Random Forest:")
print(importance_df_rf.head(10))

# Plot
plt.figure(figsize=(10, 6))
plt.barh(importance_df_rf['Feature'].head(10), importance_df_rf['Importance']
plt.xlabel('Feature Importance')
plt.title('Top 10 Most Important Features - Random Forest')
```

```
plt.tight_layout()
plt.show()
```

Top 10 Most Important Features – Random Forest:

	Feature	Importance
20	worst radius	0.145214
23	worst area	0.106068
22	worst perimeter	0.101685
7	mean concave points	0.095882
27	worst concave points	0.093164
26	worst concavity	0.087323
13	area error	0.081751
6	mean concavity	0.042094
0	mean radius	0.041257
2	mean perimeter	0.026013



In []:

```
In [36]: # Get feature importances from the best model
lr_best = lr.best_estimator_
lr_clf = lr_best.named_steps["lr"]

# Coefficients -> importance (absolute value)
coef = lr_clf.coef_ # shape: (1, n_features) for binary classification
imp = np.abs(coef).flatten()


# Create DataFrame
lr_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': imp
}).sort_values('Importance', ascending=False)

print("\nTop 10 Most Important Features – Logistic Regression:")
print(lr_importance_df.head(10))
```

```
# Plot
plt.figure(figsize=(10, 6))
plt.barh(lr_importance_df['Feature'].head(10), lr_importance_df['Importance'])
plt.xlabel('Feature Importance (Absolute Coefficient)')
plt.title('Top 10 Most Important Features – Logistic Regression')
plt.tight_layout()
plt.show()
```

Top 10 Most Important Features – Logistic Regression:

	Feature	Importance
10	radius error	1.034049
21	worst texture	0.963975
13	area error	0.861386
20	worst radius	0.858908
23	worst area	0.852374
24	worst smoothness	0.812187
7	mean concave points	0.758627
26	worst concavity	0.758144
22	worst perimeter	0.715426
27	worst concave points	0.688695



In []:

```
In [37]: # Get feature importances from the best model
svm_best = svm.best_estimator_
svm_clf = svm_best.named_steps["svm"]

# Coefficients -> importance (absolute value)
svm_coef = svm_clf.coef_ # shape: (1, n_features) for binary cl
svm_imp = np.abs(svm_coef).flatten()

# Create DataFrame
svm_importance_df = pd.DataFrame({
    'Feature': X.columns,
```

```

'Importance': svm_imp
}).sort_values('Importance', ascending=False)

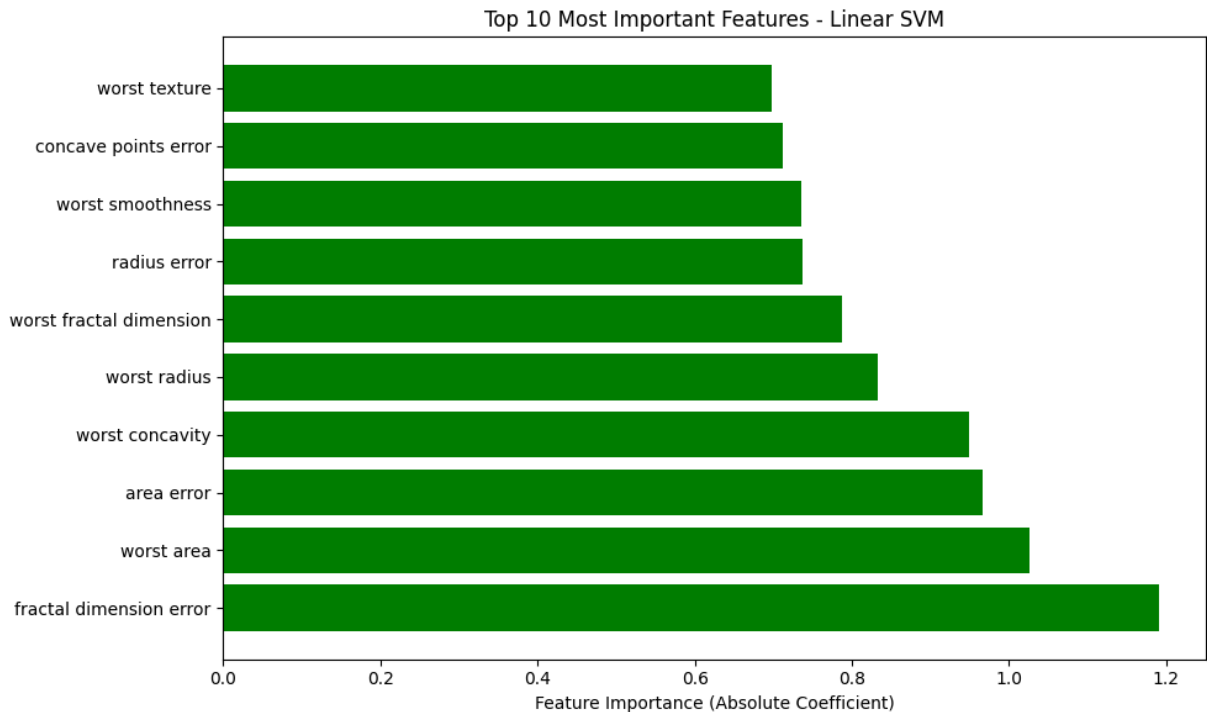
print("\nTop 10 Most Important Features - Linear SVM:")
print(svm_importance_df.head(10))

# Plot
plt.figure(figsize=(10, 6))
plt.barh(svm_importance_df['Feature'].head(10), svm_importance_df['Importance'].head(10))
plt.xlabel('Feature Importance (Absolute Coefficient)')
plt.title('Top 10 Most Important Features - Linear SVM')
plt.tight_layout()
plt.show()

```

Top 10 Most Important Features - Linear SVM:

	Feature	Importance
19	fractal dimension error	1.190984
23	worst area	1.026678
13	area error	0.966771
26	worst concavity	0.948895
20	worst radius	0.832388
29	worst fractal dimension	0.787419
10	radius error	0.736903
24	worst smoothness	0.735840
17	concave points error	0.712115
21	worst texture	0.698103



```

In [38]: # Clinical Interpretation:
          # - Tumor SIZE (radius, area, perimeter) = strongest predictor

```

```

In [ ]:

```

```

In [ ]:

```

In []: