

titanic

August 28, 2025

```
[30]: from sklearn.datasets import fetch_openml

# Load Titanic dataset from OpenML
titanic = fetch_openml(name="titanic", version=1, as_frame=True, parser="auto")
df_original = titanic.frame
print(df_original.head())
print(df_original.info())
print(f"Rows: {len(df_original)}")
```

	pclass	survived	name	sex	\
0	1	1	Allen, Miss. Elisabeth Walton	female	
1	1	1	Allison, Master. Hudson Trevor	male	
2	1	0	Allison, Miss. Helen Loraine	female	
3	1	0	Allison, Mr. Hudson Joshua Creighton	male	
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	

	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body	\
0	29.0000	0	0	24160	211.3375	B5	S	2	NaN	
1	0.9167	1	2	113781	151.5500	C22 C26	S	11	NaN	
2	2.0000	1	2	113781	151.5500	C22 C26	S	NaN	NaN	
3	30.0000	1	2	113781	151.5500	C22 C26	S	NaN	135.0	
4	25.0000	1	2	113781	151.5500	C22 C26	S	NaN	NaN	

```
home.dest
0      St Louis, MO
1  Montreal, PQ / Chesterville, ON
2  Montreal, PQ / Chesterville, ON
3  Montreal, PQ / Chesterville, ON
4  Montreal, PQ / Chesterville, ON
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1309 entries, 0 to 1308
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	pclass	1309 non-null	int64
1	survived	1309 non-null	category
2	name	1309 non-null	object
3	sex	1309 non-null	category

```

4   age          1046 non-null   float64
5   sibsp        1309 non-null   int64
6   parch        1309 non-null   int64
7   ticket       1309 non-null   object
8   fare         1308 non-null   float64
9   cabin        295 non-null    object
10  embarked     1307 non-null   category
11  boat         486 non-null    object
12  body         121 non-null    float64
13  home.dest    745 non-null    object
dtypes: category(3), float64(3), int64(3), object(5)
memory usage: 116.8+ KB
None
Rows: 1309

```

```

[31]: #1. Remove Duplicates
      # Drop duplicate rows if any
      df = df_original.drop_duplicates()
      print("Any duplicates?", df.duplicated().any())
      print("Number of duplicates:", df.duplicated().sum())
      print(f"Rows: {len(df)}")

```

```

Any duplicates? False
Number of duplicates: 0
Rows: 1309

```

```

[32]: #2. Handle Missing Values

      # Check missing values
      print(df.isnull().sum())

```

```

pclass          0
survived         0
name            0
sex             0
age            263
sibsp           0
parch           0
ticket          0
fare            1
cabin          1014
embarked         2
boat           823
body           1188
home.dest       564
dtype: int64

```

```
[33]: # fill with mean/median:
df["age"].fillna(df["age"].median(), inplace=True)
df["fare"].fillna(df["fare"].median(), inplace=True)
```

```
[34]: # Only 2 missing → fill with the mode (most common port).
df["embarked"].fillna(df["embarked"].mode()[0], inplace=True)
```

```
[35]: # Drop cabin, boat, body, home.dest:
# Since they are >70% missing, usually dropped.
df.drop(columns=["cabin", "boat", "body", "home.dest"], inplace=True)
```

```
[36]: df
```

```
[36]:      pclass  survived      name \
0         1         1  Allen, Miss. Elisabeth Walton
1         1         1  Allison, Master. Hudson Trevor
2         1         0  Allison, Miss. Helen Loraine
3         1         0  Allison, Mr. Hudson Joshua Creighton
4         1         0  Allison, Mrs. Hudson J C (Bessie Waldo Daniels)
...     ...     ...
1304      3         0  Zabour, Miss. Hileni
1305      3         0  Zabour, Miss. Thamine
1306      3         0  Zakarian, Mr. Mapriededer
1307      3         0  Zakarian, Mr. Ortin
1308      3         0  Zimmerman, Mr. Leo

      sex    age  sibsp  parch  ticket    fare  embarked
0  female  29.0000     0     0   24160  211.3375        S
1   male    0.9167     1     2  113781  151.5500        S
2  female    2.0000     1     2  113781  151.5500        S
3   male   30.0000     1     2  113781  151.5500        S
4  female   25.0000     1     2  113781  151.5500        S
...     ...     ...     ...     ...     ...
1304 female   14.5000     1     0   2665   14.4542        C
1305 female   28.0000     1     0   2665   14.4542        C
1306  male   26.5000     0     0   2656    7.2250        C
1307  male   27.0000     0     0   2670    7.2250        C
1308  male   29.0000     0     0  315082    7.8750        S
```

[1309 rows x 10 columns]

```
[9]: # Check missing values
print(df.isnull().sum())
```

```
pclass    0
survived   0
name       0
sex        0
```

```

age          0
sibsp        0
parch        0
ticket       0
fare         0
embarked     0
dtype: int64

```

[37]: #3. Check inconsistent formats

```
print(df.isin(["N/A", "na", "NA", "null", "NULL", "?", "--"]).sum())
```

```

pclass       0
survived     0
name         0
sex          0
age          0
sibsp        0
parch        0
ticket       0
fare         0
embarked     0
dtype: int64

```

[38]: # check unique values:

```
print(df['sex'].unique())
```

```

['female', 'male']
Categories (2, object): ['female', 'male']

```

[39]: #4. Check Outlier with histogram

```
import matplotlib.pyplot as plt
```

```
# selects only the numeric columns from the DataFrame df.
```

```
numeric_cols = df.select_dtypes(include=["int64", "float64"]).columns
```

```
for col in numeric_cols:
```

```
    plt.figure(figsize=(6,4)) # create a new figure width = 6 inches, height = 4 inches
```

```
    plt.hist(df[col].dropna(), bins=30, color="teal", edgecolor="black")
```



```
    #divides the range into 30 equal-width "buckets."
```



```
    plt.title(f"Histogram of {col}")
```


```
    plt.xlabel(col)
```

```
    plt.ylabel("Frequency")
```

```
    plt.show()
```







```
[27]: # 5. check Data Types  
print(df.dtypes)
```

```
pclass      int64  
survived     category  
name         object  
sex          category  
age          float64  
sibsp        int64  
parch        int64  
ticket       object  
fare         float64  
embarked     category  
dtype: object
```

```
[28]: df['survived'] = df['survived'].astype('int64')    # keep 0/1
```

```
[ ]:
```